

GCP Dual-Stack VPN Solution

HA VPN with IPv4 and IPv6 BGP Sessions

Technical Implementation Guide

January 2026

Contents

1	Executive Summary	2
2	Architecture Overview	3
2.1	Network Topology	3
2.2	Key Components	3
2.3	IPv6 Addressing Strategy	4
3	Critical Design Decisions	4
3.1	Why Dedicated IPv6 BGP Sessions?	4
3.2	IPv6 Address Range Validity	4
3.3	GRE Tunnel for Custom Prefixes	5
3.4	BGP Session Configuration	5
4	Terraform Implementation	5
4.1	Project Structure	5
4.2	Key Resource Definitions	6
4.2.1	GCP VPC with IPv6	6
4.2.2	HA VPN Gateway (Dual-Stack)	6
4.2.3	IPv6 BGP Session (The Key!)	6
4.2.4	On-Prem Router Configuration (LibreSwan)	7
4.2.5	On-Prem Router Configuration (FRR BGP)	7
4.2.6	Firewall Rules (Separate IPv4/IPv6)	7
5	Deployment	8
5.1	Prerequisites	8
5.2	Deployment Steps	8
5.3	Deployment Time	8
6	Verification	9
6.1	Check IPsec Tunnel Status (On-Prem)	9
6.2	Check BGP Session Status	9
6.3	Verify Route Installation	9
6.4	Test Connectivity	10
7	Troubleshooting	10
7.1	Common Issues	10
7.2	LibreSwan Dual-Stack Issue	10
7.3	Diagnostic Commands	11
8	Cloud Provider Comparison	11
9	Summary	11
10	Cleanup	11

1 Executive Summary

This document describes the implementation of a dual-stack (IPv4 + IPv6) Site-to-Site VPN solution using Google Cloud Platform's HA VPN service connecting to an on-premises simulation environment using LibreSwan IPsec and FRR BGP routing.


Capability	IPv4	IPv6
VPN Tunnels	✓	✓
BGP Sessions	✓	✓
Route Exchange	✓	✓
Cross-VPN Ping	38ms	33-37ms
Custom Prefix Advertisement	✓	✓

Key Achievements:

- Full dual-stack connectivity with automatic route propagation via dedicated IPv4 and IPv6 BGP sessions
- On-prem simulation using LibreSwan + FRR successfully peers with GCP Cloud Router
- Custom IPv6 prefix advertisement (2001:db8:beef::/48) via BGP

2 Architecture Overview

2.1 Network Topology



Listing 1: Network Architecture with On-Prem Simulation

2.2 Key Components

Component	Description
-----------	-------------

HA VPN Gateway	Regional resource with <code>stack_type = "IPV4_IPV6"</code> for dual-stack tunnel support
External VPN Gateway	Represents on-prem router public IPs with <code>redundancy_type = "TWO_IPS_REDUNDANCY"</code>
Cloud Router	BGP router with separate IPv4 and IPv6 peering sessions (ASN 65515)
LibreSwan	IPsec VPN daemon on on-prem routers with VTI interfaces
FRR	Free Range Routing daemon for BGP on on-prem routers (ASN 65001)
GRE Tunnel	Carries traffic for custom-advertised prefixes not in VPC range

2.3 IPv6 Addressing Strategy

Component	IPv6 Range	Purpose
GCP Cloud VPC	fd20:f:1::/48	ULA for GCP internal
On-prem VPC (internal)	fd20:e:1::/48	ULA for on-prem internal
On-prem BGP advertisement	2001:db8:beef::/48	Custom prefix advertised via BGP
BGP peering	fdff:1::/64	Link-local for BGP sessions

3 Critical Design Decisions

3.1 Why Dedicated IPv6 BGP Sessions?

Important: MP-BGP (Multiprotocol BGP) on IPv4 sessions does NOT properly install IPv6 routes in GCP Cloud Router.

What doesn't work:

- Setting `enable_ipv6 = true` on IPv4 BGP peers
- IPv6 routes are advertised but NOT installed into VPC routing table
- Routes appear in `bestRoutes` but with IPv4 next-hops (unusable for IPv6 forwarding)

What works:

- Dedicated IPv6 BGP sessions with IPv6 peering addresses
- Using the reserved `fdff:1::/64` ULA range for BGP peering
- /126 subnets for point-to-point links

3.2 IPv6 Address Range Validity

Critical Finding: The Linux kernel will NOT route to IPv6 addresses in unallocated ranges. `dead:beef::/48` does NOT work because it's outside valid IPv6 address space.

Valid IPv6 ranges for routing:

Range	Type	Routable
2000::/3	Global Unicast (starts with 2 or 3)	✓ Yes

fc00::/7	Unique Local (starts with fc or fd)	✓ Yes
dead::/16	Unallocated (starts with d)	✗ No
face::/16	Unallocated	✗ No

Why `dead:beef::/48` fails:


- dead in hex = 1101 1110 1010 1101 in binary
- First 3 bits are 110, not 001 (required for global unicast)
- Linux kernel refuses to use routes for these addresses
- Kernel falls back to default gateway instead of specific route

Working alternative:

- Use `2001:db8:beef::/48` instead (documentation range, valid global unicast)
- Or use `fdbe:ef00::/48` (ULA range, looks like “beef”)

3.3 GRE Tunnel for Custom Prefixes

Since custom-advertised prefixes (like `2001:db8:beef::/48`) are not part of the on-prem VPC’s allocated range, a GRE tunnel is needed to route traffic to the test VM:



The GRE tunnel uses the VPC’s internal ULA range for transport, and the test VM has the custom prefix configured on the tunnel interface.

3.4 BGP Session Configuration

Session	GCP Cloud IP	On-Prem IP	Routes
IPv4 Tunnel 0	169.254.0.1/30	169.254.0.2/30	IPv4 only
IPv4 Tunnel 1	169.254.1.1/30	169.254.1.2/30	IPv4 only
IPv6 Tunnel 0	fdff:1::1/126	fdff:1::2/126	IPv6 only
IPv6 Tunnel 1	fdff:1::5/126	fdff:1::6/126	IPv6 only

Table 1: BGP Peering Address Allocation

4 Terraform Implementation

4.1 Project Structure

```

cloud_provi/
├── terraform-gcp-external/      # GCP Cloud Infrastructure
│   ├── main.tf                  # VPC, subnets, firewall
│   ├── vpn.tf                   # HA VPN Gateway, tunnels
│   ├── bgp.tf                   # Cloud Router, BGP peers
│   ├── test-vm.tf               # Test VM
│   ├── variables.tf             # Input variables
│   └── outputs.tf               # VPN config for on-prem
└── terraform-onprem-sim/       # On-Prem Simulation
    ├── main.tf                  # VPC, routers, test VM
    ├── cloud-init-router.yaml   # LibreSwan + FRR bootstrap
    └── scripts/
        └── configure-gcp-vpn.sh # GCP VPN config script
  
```

4.2 Key Resource Definitions

4.2.1 GCP VPC with IPv6

```
resource "google_compute_network" "cloud" {
  name          = "vpc-gcp-cloud-prod"
  auto_create_subnetworks = false
  routing_mode      = "GLOBAL"
  enable_ula_internal_ipv6 = true
  internal_ipv6_range = "fd20:f:1::/48"
  delete_default_routes_on_create = false
}

resource "google_compute_subnetwork" "cloud_workload" {
  name          = "subnet-workload-prod"
  ip_cidr_range = "10.10.1.0/24"
  region        = "us-central1"
  network       = google_compute_network.cloud.id
  stack_type    = "IPV4_IPV6"
  ipv6_access_type = "INTERNAL"
}
```

4.2.2 HA VPN Gateway (Dual-Stack)

```
resource "google_compute_ha_vpn_gateway" "cloud" {
  name          = "vpngw-cloud-prod"
  region        = "us-central1"
  network       = google_compute_network.cloud.id
  stack_type    = "IPV4_IPV6" # Critical for dual-stack
}

resource "google_compute_external_vpn_gateway" "onprem" {
  name          = "extgw-onprem-prod"
  redundancy_type = "TWO_IPS_REDUNDANCY"

  interface {
    id          = 0
    ip_address = var.onprem_router_1_ip # Router 1 public IP
  }
  interface {
    id          = 1
    ip_address = var.onprem_router_2_ip # Router 2 public IP
  }
}
```

4.2.3 IPv6 BGP Session (The Key!)

```
resource "google_compute_router_interface" "cloud_interface_0_v6" {
  name          = "interface-cloud-0-v6"
  router        = google_compute_router.cloud.name
  region        = "us-central1"
  ip_range     = "fdff:1::1/126" # IPv6 ULA address
  vpn_tunnel   = google_compute_vpn_tunnel.cloud_to_onprem_0.name
}

resource "google_compute_router_peer" "cloud_peer_0_v6" {
  name          = "peer-cloud-to-onprem-0-v6"
  router        = google_compute_router.cloud.name
  region        = "us-central1"
  peer_ip_address = "fdff:1::2" # On-prem router IPv6
  peer_asn      = 65001
```

```

interface      = google_compute_router_interface.cloud_interface_0_v6.name
enable_ipv6   = true
}

```

4.2.4 On-Prem Router Configuration (LibreSwan)

```

# /etc/ipsec.d/gcp-tunnel-r1.conf
conn gcp-tunnel-r1
    authby=secret
    auto=start
    left=%defaultroute
    leftid=34.82.67.66
    right=35.242.119.182
    type=tunnel
    ikev2=yes
    ike=aes256-sha256-modp2048
    esp=aes256-sha256
    ikelifetime=36000s
    salifetime=10800s
    leftsubnet=0.0.0.0/0
    rightsubnet=0.0.0.0/0
    mark=200/0xffffffff
    vti-interface=vti20
    vti-routing=no

```

4.2.5 On-Prem Router Configuration (FRR BGP)

```

router bgp 65001
    bgp router-id 192.168.0.10
    no bgp ebgp-requires-policy
    no bgp network import-check

    ! IPv4 neighbor
    neighbor 169.254.0.1 remote-as 65515
    neighbor 169.254.0.1 ebgp-multipath 2

    ! IPv6 neighbor (dedicated session)
    neighbor fdff:1::1 remote-as 65515
    neighbor fdff:1::1 ebgp-multipath 2

    address-family ipv4 unicast
        network 192.168.0.0/16
        neighbor 169.254.0.1 activate
    exit-address-family

    address-family ipv6 unicast
        network fd20:e:1::/48
        network 2001:db8:beef::/48
        neighbor fdff:1::1 activate
    exit-address-family
exit

```

4.2.6 Firewall Rules (Separate IPv4/IPv6)

```

# IPv4 firewall rule
resource "google_compute_firewall" "cloud_allow_internal" {
    name      = "fw-cloud-allow-internal"
    network  = google_compute_network.cloud.id
    allow { protocol = "icmp" }
    allow { protocol = "tcp" }
    allow { protocol = "udp" }

```

```

    source_ranges = ["10.10.0.0/16", "192.168.0.0/16"]
}

# IPv6 firewall rule (MUST be separate)
resource "google_compute_firewall" "cloud_allow_internal_ipv6" {
  name      = "fw-cloud-allow-internal-ipv6"
  network   = google_compute_network.cloud.id
  allow { protocol = "58" } # ICMPv6
  allow { protocol = "tcp" }
  allow { protocol = "udp" }
  source_ranges = ["fd20:e:1::/48", "fd20:f:1::/48", "2001:db8:beef::/48"]
}

```

GCP Limitation: Firewall rules cannot mix IPv4 and IPv6 in the same `source_ranges`. Create separate rules for each address family.

5 Deployment

5.1 Prerequisites

- GCP Project with billing enabled
- Compute Engine API enabled
- Terraform >= 1.5.0
- gcloud CLI authenticated

5.2 Deployment Steps

```

# Step 1: Deploy On-Prem Simulation
cd terraform-onprem-sim
terraform init
terraform apply -var="project_id=your-project-id"

# Note the router public IPs from output

# Step 2: Deploy GCP Cloud Infrastructure
cd ../../terraform-gcp-external
terraform apply \
  -var="project_id=your-project-id" \
  -var="onprem_router_1_ip=<router-1-ip>" \
  -var="onprem_router_2_ip=<router-2-ip>"

# Step 3: Configure On-Prem Routers
# Save config from terraform output
terraform output -raw onprem_router_1_json_config > /tmp/r1.json

# SSH to Router 1 and configure
gcloud compute ssh vm-router-1-prod --zone=us-west1-a
sudo /usr/local/bin/configure-gcp-vpn.sh --config /tmp/r1.json --router-id 1

# Repeat for Router 2

```

5.3 Deployment Time

Resource	Time
On-Prem VPC + Routers	2 minutes
GCP VPC + VPN Gateway	1 minute

VPN Tunnels + BGP	30 seconds
Router Configuration	1 minute
Total	5 minutes

6 Verification

6.1 Check IPsec Tunnel Status (On-Prem)

```
gcloud compute ssh vm-router-1-prod --zone=us-west1-a
sudo ipsec status
```

Expected Output:

```
gcp-tunnel-r1[1]: ESTABLISHED 1 hour ago
gcp-tunnel-r1{1}: INSTALLED, TUNNEL
```

6.2 Check BGP Session Status

```
# On-prem router
sudo vtysh -c "show bgp summary"

# GCP Cloud Router
gcloud compute routers get-status router-cloud-prod-us-central1 \
--region=us-central1 \
--format="yaml(result.bgpPeerStatus[].name,result.bgpPeerStatus[].state)"
```

Expected Output:

```
result:
  bgpPeerStatus:
    - name: peer-to-onprem-0-v4
      state: Established
    - name: peer-to-onprem-0-v6
      state: Established
    - name: peer-to-onprem-1-v4
      state: Established
    - name: peer-to-onprem-1-v6
      state: Established
```

6.3 Verify Route Installation

```
gcloud compute routers get-status router-cloud-prod-us-central1 \
--region=us-central1 \
--format=json | jq '.result.bestRoutesForRouter[] |
  select(.destRange | contains("2001:db8") or contains("fd20:e"))'
```

Expected Output:

```
{
  "destRange": "2001:db8:beef::/48",
  "nextHopIp": "fdff:1::2",
  "routeStatus": "ACTIVE",
  "routeType": "BGP"
}
{
  "destRange": "fd20:e:1::/48",
  "nextHopIp": "fdff:1::2",
  "routeStatus": "ACTIVE",
  "routeType": "BGP"
}
```

6.4 Test Connectivity

```
# SSH to GCP cloud VM
gcloud compute ssh vm-cloud-test-prod-us-central1 --zone=us-central1-a

# Test IPv4 to on-prem
ping -c 3 192.168.1.100

# Test IPv6 to on-prem ULA
ping6 -c 3 fd20:e:1:2001:0:1::

# Test IPv6 to custom-advertised prefix
ping6 -c 3 2001:db8:beef::100
```

Expected Results:

```
PING 192.168.1.100: 3 packets, 0% loss, rtt avg 38ms
PING fd20:e:1:2001:0:1:: 3 packets, 0% loss, rtt avg 37ms
PING 2001:db8:beef::100: 3 packets, 0% loss, rtt avg 33ms
```

7 Troubleshooting

7.1 Common Issues

Symptom	Solution
IPv6 BGP session DOWN	Verify /126 subnets don't overlap. Check VTI interface has IPv6 address.
IPv6 routes not installed	Don't use MP-BGP on IPv4 sessions. Create dedicated IPv6 BGP sessions.
Firewall blocking traffic	Create separate IPv4 and IPv6 firewall rules.
Ping6 to custom prefix fails	Ensure prefix is in valid range (2000::/3 or fd00::/8). Use GRE tunnel.
dead:beef:: not routable	Invalid IPv6 range. Use 2001:db8:beef:: instead.
LibreSwan crash with dual-stack	Create separate IPv4 and IPv6 IPsec connections, not combined.

7.2 LibreSwan Dual-Stack Issue

LibreSwan 3.x Limitation: Using `leftsubnet=0.0.0.0/0, ::/0` (combined IPv4+IPv6 traffic selectors) can cause daemon crashes.

Workaround: Create separate IPsec connections for IPv4 and IPv6, both using the same VTI interface and mark:

```
# IPv4 connection
conn gcp-tunnel-r1-v4
  leftsubnet=0.0.0.0/0
  rightsubnet=0.0.0.0/0
  mark=200/0xffffffff
  vti-interface=vti20

# IPv6 connection (same VTI)
conn gcp-tunnel-r1-v6
  leftsubnet=::/0
  rightsubnet=::/0
```

```
mark=200/0xffffffff
vti-interface=vti20
```

7.3 Diagnostic Commands

```
# On-prem router
sudo ipsec status          # IPsec tunnel status
sudo vtysh -c "show bgp summary" # BGP peer status
sudo vtysh -c "show bgp ipv6"   # IPv6 routes
ip -6 route show            # Kernel IPv6 routes
ip -6 route get 2001:db8:beef::100 # Test route lookup

# GCP
gcloud compute routers get-status <router> --region=<region>
gcloud compute vpn-tunnels describe <tunnel> --region=<region>
gcloud compute routes list --filter="network:<vpc>"
```

8 Cloud Provider Comparison

Capability	Azure VWAN	Azure VPN GW	GCP HA VPN
IPv6 VPN config	×	Preview	✓ GA
IPv6 BGP sessions	×	Preview	✓ GA
IPv6 route learning	×	Preview	✓ GA
Cross-VPN IPv4	✓	✓ 69ms	✓ 38ms
Cross-VPN IPv6	×	Untested	✓ 33ms
Deployment time	30-45 min	30-45 min	3-5 min
Custom IPv6 prefix via BGP	×	×	✓

Conclusion: GCP HA VPN provides full dual-stack support with the fastest deployment time. Azure requires preview enrollment for IPv6 VPN support.

9 Summary

Working Configuration:

- GCP HA VPN with `stack_type = "IPV4_IPV6"`
- Dedicated IPv6 BGP sessions (not MP-BGP on IPv4)
- LibreSwan on-prem with separate IPv4/IPv6 IPsec connections
- FRR BGP with dual address families
- GRE tunnel for routing custom-advertised prefixes

Key Findings:

1. MP-BGP over IPv4 sessions does not install IPv6 routes in GCP
2. `dead:beef::/48` fails because it's in unallocated IPv6 space
3. Use valid prefixes: `2001:db8::/32` (documentation) or `fd00::/8` (ULA)
4. LibreSwan requires separate connections for IPv4 and IPv6 traffic selectors

10 Cleanup

To destroy all resources:

```
# Destroy GCP cloud infrastructure
cd terraform-gcp-external
terraform destroy

# Destroy on-prem simulation
```

```
cd ../../terraform-onprem-sim  
terraform destroy
```

Document generated from successful GCP dual-stack VPN implementation

Project: dual-stack-vpn-test | Regions: us-central1 (GCP), us-west1 (On-Prem Sim)

Test Results: IPv4 38ms, IPv6 ULA 37ms, IPv6 Custom Prefix 33ms | 0% packet loss