

GCP Dual-Stack VPN Solution

HA VPN with IPv4 and IPv6 BGP Sessions

Technical Implementation Guide

January 2026

Contents

| | | |
|-------|---|---|
| 1 | Executive Summary | 2 |
| 2 | Architecture Overview | 2 |
| 2.1 | Network Topology | 2 |
| 2.2 | Key Components | 2 |
| 3 | Critical Design Decision | 3 |
| 3.1 | Why Dedicated IPv6 BGP Sessions? | 3 |
| 3.2 | BGP Session Configuration | 3 |
| 4 | Terraform Implementation | 3 |
| 4.1 | Project Structure | 3 |
| 4.2 | Key Resource Definitions | 3 |
| 4.2.1 | Local Variables | 3 |
| 4.2.2 | VPC Networks with IPv6 | 4 |
| 4.2.3 | HA VPN Gateway (Dual-Stack) | 4 |
| 4.2.4 | IPv4 BGP Session | 4 |
| 4.2.5 | IPv6 BGP Session (The Key!) | 5 |
| 4.2.6 | Firewall Rules (Separate IPv4/IPv6) | 5 |
| 5 | Deployment | 5 |
| 5.1 | Prerequisites | 5 |
| 5.2 | Deployment Commands | 6 |
| 5.3 | Deployment Time | 6 |
| 6 | Verification | 6 |
| 6.1 | Check BGP Session Status | 6 |
| 6.2 | Verify Route Installation | 6 |
| 6.3 | Test Connectivity | 7 |
| 7 | Troubleshooting | 7 |
| 7.1 | Common Issues | 7 |
| 7.2 | Diagnostic Commands | 7 |
| 8 | Comparison: Azure vs GCP | 7 |
| 9 | Cleanup | 8 |

1 Executive Summary

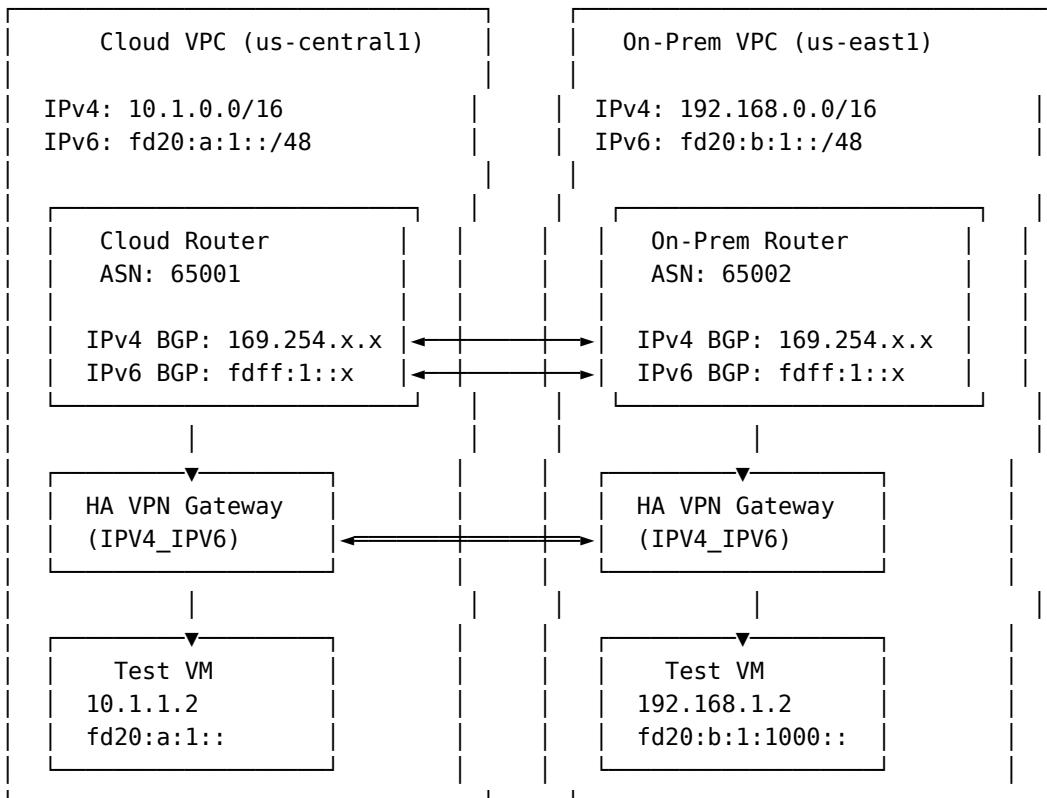
This document describes the implementation of a dual-stack (IPv4 + IPv6) Site-to-Site VPN solution using Google Cloud Platform's HA VPN service. The solution enables full IPv4 and IPv6 connectivity between two VPC networks using BGP dynamic routing.

| Capability | IPv4 | IPv6 |
|----------------|------|------|
| VPN Tunnels | ✓ | ✓ |
| BGP Sessions | ✓ | ✓ |
| Route Exchange | ✓ | ✓ |
| Cross-VPN Ping | 32ms | 33ms |

Key Achievement: Full dual-stack connectivity with automatic route propagation via dedicated IPv4 and IPv6 BGP sessions.

2 Architecture Overview

2.1 Network Topology



Listing 1: Network Architecture Diagram

2.2 Key Components

| Component | Description |
|----------------|--|
| HA VPN Gateway | Regional resource with <code>stack_type = "IPV4_IPV6"</code> for dual-stack tunnel support |
| Cloud Router | BGP router with separate IPv4 and IPv6 peering sessions |
| VPN Tunnels | 4 tunnels total (2 per direction) for 99.99% SLA |
| BGP Sessions | 8 total: 4 IPv4 sessions + 4 IPv6 sessions |

3 Critical Design Decision

3.1 Why Dedicated IPv6 BGP Sessions?

Important: MP-BGP (Multiprotocol BGP) on IPv4 sessions does NOT work for IPv6 route installation when using GCP Cloud Routers on both ends.

Why this is GCP-specific:

GCP Cloud Router does not allow manual manipulation of the BGP next-hop attribute. When MP-BGP advertises IPv6 routes over an IPv4 session, the next-hop remains the IPv4 peering address (e.g., 169.254.x.x), which cannot be used to forward IPv6 traffic.

With traditional routers (Cisco, Juniper, etc.), you can use MP-BGP with IPv4 transport and manually set the IPv6 next-hop using route-maps or policies. This allows IPv6 routes to be advertised with a valid IPv6 next-hop address, making them installable.

What doesn't work (GCP-to-GCP):

- Setting `enable_ipv6 = true` on IPv4 BGP peers
- IPv6 routes are advertised but NOT installed into VPC routing table
- Routes appear in `bestRoutes` but with IPv4 next-hops (unusable)

What works (GCP-to-GCP):

- Dedicated IPv6 BGP sessions with IPv6 peering addresses
- Using the reserved `fdff:1::/64` ULA range for BGP peering
- /126 subnets for point-to-point links

Alternative (GCP-to-physical router):

- MP-BGP on IPv4 session may work if the remote router can set the IPv6 next-hop
- Requires route-map/policy on the non-GCP side to rewrite next-hop

3.2 BGP Session Configuration

| Session | Cloud IP | On-Prem IP | Routes |
|---------------|----------------------------|----------------------------|-----------|
| IPv4 Tunnel 0 | 169.254.0.1/30 | 169.254.0.2/30 | IPv4 only |
| IPv4 Tunnel 1 | 169.254.1.1/30 | 169.254.1.2/30 | IPv4 only |
| IPv6 Tunnel 0 | <code>fdff:1::1/126</code> | <code>fdff:1::2/126</code> | IPv6 only |
| IPv6 Tunnel 1 | <code>fdff:1::5/126</code> | <code>fdff:1::6/126</code> | IPv6 only |

Table 1: BGP Peering Address Allocation

4 Terraform Implementation

4.1 Project Structure

```
terraform-gcp/
├── main.tf          # All resources
├── terraform.tfvars # Configuration values
└── README.md        # Documentation
```

4.2 Key Resource Definitions

4.2.1 Local Variables

```
locals {
  # VPC Address Ranges
  cloud_ipv4_cidr = "10.1.0.0/16"
```

```

cloud_ipv6_cidr = "fd20:a:1::/48"
onprem_ipv4_cidr = "192.168.0.0/16"
onprem_ipv6_cidr = "fd20:b:1::/48"

# BGP ASNs
cloud_asn = 65001
onprem_asn = 65002

# IPv6 BGP Peering Addresses (fdff:1::/64 ULA range)
bgp_v6_cloud_0 = "fdff:1::1/126"
bgp_v6_onprem_0 = "fdff:1::2/126"
bgp_v6_cloud_1 = "fdff:1::5/126"
bgp_v6_onprem_1 = "fdff:1::6/126"
}

```

4.2.2 VPC Networks with IPv6

```

resource "google_compute_network" "cloud" {
  name           = "vpc-cloud-prod"
  auto_create_subnetworks = false
  routing_mode      = "GLOBAL"
  enable_ula_internal_ipv6 = true
  internal_ipv6_range = "fd20:a:1::/48"
}

resource "google_compute_subnetwork" "cloud_workload" {
  name           = "subnet-workload-prod"
  ip_cidr_range   = "10.1.1.0/24"
  region          = "us-central1"
  network         = google_compute_network.cloud.id
  stack_type      = "IPV4_IPV6"
  ipv6_access_type = "INTERNAL"
}

```

4.2.3 HA VPN Gateway (Dual-Stack)

```

resource "google_compute_ha_vpn_gateway" "cloud" {
  name       = "vpngw-cloud-prod"
  region     = "us-central1"
  network    = google_compute_network.cloud.id
  stack_type = "IPV4_IPV6" # Critical for dual-stack
}

```

4.2.4 IPv4 BGP Session

```

resource "google_compute_router_interface" "cloud_interface_0" {
  name       = "interface-cloud-0"
  router     = google_compute_router.cloud.name
  region     = "us-central1"
  ip_range   = "169.254.0.1/30"
  vpn_tunnel = google_compute_vpn_tunnel.cloud_to_onprem_0.name
}

resource "google_compute_router_peer" "cloud_peer_0" {
  name       = "peer-cloud-to-onprem-0"
  router     = google_compute_router.cloud.name
  region     = "us-central1"
  peer_ip_address = "169.254.0.2"
  peer_asn    = 65002
  interface   = google_compute_router_interface.cloud_interface_0.name
}

```

```

# Note: NO enable_ipv6 here - IPv4 only
}



### 4.2.5 IPv6 BGP Session (The Key!)


resource "google_compute_router_interface" "cloud_interface_0_v6" {
  name        = "interface-cloud-0-v6"
  router      = google_compute_router.cloud.name
  region      = "us-central1"
  ip_range    = "fdff:1::1/126" # IPv6 ULA address
  vpn_tunnel = google_compute_vpn_tunnel.cloud_to_onprem_0.name
}

resource "google_compute_router_peer" "cloud_peer_0_v6" {
  name        = "peer-cloud-to-onprem-0-v6"
  router      = google_compute_router.cloud.name
  region      = "us-central1"
  peer_ip_address = "fdff:1::2" # IPv6 peer address
  peer_asn     = 65002
  interface    = google_compute_router_interface.cloud_interface_0_v6.name
  enable_ipv6  = true # Enable IPv6 route exchange
}

```

4.2.6 Firewall Rules (Separate IPv4/IPv6)

```

# IPv4 firewall rule
resource "google_compute_firewall" "cloud_allow_internal" {
  name        = "fw-cloud-allow-internal"
  network    = google_compute_network.cloud.id
  allow { protocol = "icmp" }
  allow { protocol = "tcp" }
  allow { protocol = "udp" }
  source_ranges = ["10.1.0.0/16", "192.168.0.0/16"]
}

# IPv6 firewall rule (MUST be separate)
resource "google_compute_firewall" "cloud_allow_internal_ipv6" {
  name        = "fw-cloud-allow-internal-ipv6"
  network    = google_compute_network.cloud.id
  allow { protocol = "58" } # ICMPv6
  allow { protocol = "tcp" }
  allow { protocol = "udp" }
  source_ranges = ["fd20:a:1::/48", "fd20:b:1::/48"]
}

```

GCP Limitation: Firewall rules cannot mix IPv4 and IPv6 in the same `source_ranges`. Create separate rules for each address family.

5 Deployment

5.1 Prerequisites

- GCP Project with billing enabled
- Compute Engine API enabled
- Terraform >= 1.5.0
- gcloud CLI authenticated

5.2 Deployment Commands

```
# Initialize Terraform
cd terraform-gcp
terraform init

# Create terraform.tfvars
cat > terraform.tfvars << EOF
project_id      = "your-project-id"
region          = "us-central1"
onprem_region   = "us-east1"
environment     = "prod"
vpn_shared_secret = "YourSecureSharedSecret123!"
EOF

# Plan and apply
terraform plan -out=tfplan
terraform apply tfplan
```

5.3 Deployment Time

| Resource | Time |
|-----------------|------------------|
| VPC Networks | 20 seconds |
| HA VPN Gateways | 15 seconds |
| VPN Tunnels | 15 seconds |
| BGP Sessions | 30 seconds |
| Total | 3 minutes |

6 Verification

6.1 Check BGP Session Status

```
gcloud compute routers get-status router-cloud-prod \
--region=us-central1 \
--format="yaml(result.bgpPeerStatus[].name,result.bgpPeerStatus[].state)"
```

Expected Output:

```
result:
  bgpPeerStatus:
  - name: peer-cloud-to-onprem-0
    state: Established
  - name: peer-cloud-to-onprem-0-v6
    state: Established
  - name: peer-cloud-to-onprem-1
    state: Established
  - name: peer-cloud-to-onprem-1-v6
    state: Established
```

6.2 Verify Route Installation

```
gcloud compute routers get-status router-cloud-prod \
--region=us-central1 \
--format="yaml(result.bestRoutes)" | grep -E "destRange|nextHopIp"
```

Expected Output (showing IPv6 routes with IPv6 next-hops):

```
destRange: fd20:b:1::/48
nextHopIp: ffff:1::2      # IPv6 next-hop!
```

```
destRange: fd20:b:1:1000::/64
nextHopIp: fdff:1::2
```

6.3 Test Connectivity

```
# SSH to cloud VM
gcloud compute ssh vm-cloud-test-prod --zone=us-central1-a

# Test IPv4
ping -c 3 192.168.1.2

# Test IPv6
ping6 -c 3 fd20:b:1:1000::
```

Expected Results:

```
PING 192.168.1.2: 3 packets transmitted, 3 received, 0% packet loss
rtt min/avg/max = 32.2/32.5/32.8 ms
```

```
PING fd20:b:1:1000::: 3 packets transmitted, 3 received, 0% packet loss
rtt min/avg/max = 32.5/33.3/34.7 ms
```

7 Troubleshooting

7.1 Common Issues

| Symptom | Solution |
|---------------------------|---|
| IPv6 BGP session DOWN | Verify /126 subnets don't overlap. Each tunnel needs unique /126 range. |
| IPv6 routes not installed | Don't use MP-BGP on IPv4 sessions. Create dedicated IPv6 BGP sessions. |
| Firewall blocking traffic | Create separate IPv4 and IPv6 firewall rules. Can't mix address families. |
| Ping6 100% packet loss | Check bestRoutes for IPv6 next-hops. Must be fdff:1::x, not 169.254.x.x. |

7.2 Diagnostic Commands

```
# Check all routes in router
gcloud compute routers get-status router-cloud-prod \
--region=us-central1 --format=json | jq '.result.bestRoutes[]'

# Check VPC routes
gcloud compute routes list --filter="network:vpc-cloud-prod"

# Check firewall rules
gcloud compute firewall-rules list --filter="network:vpc-cloud-prod"
```

8 Comparison: Azure vs GCP

| Capability | Azure VWAN | GCP HA VPN |
|-------------------------|--------------|------------|
| IPv6 in VPN config | ✗ | ✓ |
| IPv6 BGP sessions | ✗ | ✓ |
| IPv6 route installation | ✗ | ✓ |
| Cross-VPN IPv4 | ✗ (VM issue) | ✓ 32ms |

| | | |
|-----------------|-----------|--------|
| Cross-VPN IPv6 | ✗ | ✓ 33ms |
| Deployment time | 30-45 min | 3 min |

Conclusion: GCP HA VPN provides full dual-stack support when using dedicated IPv6 BGP sessions. Azure Virtual WAN does not support IPv6 over S2S VPN.

9 Cleanup

To destroy all resources:

```
cd terraform-gcp
terraform destroy
```

Document generated from successful GCP dual-stack VPN implementation

Project: dual-stack-vpn-test | Regions: us-central1, us-east1